

The background of the slide is a complex, abstract network of black dots (nodes) connected by thin, light gray lines (edges). The nodes are distributed across the entire frame, with some areas having higher concentrations of nodes and edges, creating a sense of depth and connectivity. The overall aesthetic is technical and modern.

Project #3: Matrix-free LSQR

Presented by Yuval Freund and Dominik Marino

Table of content

Recap: Working plan

What we have done

LSQR implementation

Why cuBLAS?

Testing

Problems

Our next steps

Working Plan

Phase 0: Understanding of LSQR method

Phase 1: Programm and test of the LSQR method with a small matrix

Phase 2: Try and modify LSQR method for a matrix **A** that can not fully stored in the memory

What we have done:

- *Solving the LSQR problem (Yuval):*
 - Look up for useful libraries
 - Implementing the LSQR problem in CUDA
 - Solving compilation & linking issues, debugging and verifying correctness
- *Testing (Dominik):*
 - Find suitable and different test matrices and vectors
 - Implement the given LSMR library for any input data
 - Compare results from the CUDA implementing and the LSMR library

LSQR implementation

The algorithm in the article

Algorithm LSQR

(1) (Initialize.)

$$\beta_1 u_1 = b, \quad \alpha_1 v_1 = A^T u_1, \quad w_1 = v_1, \quad x_0 = 0,$$

$$\bar{\phi}_1 = \beta_1, \quad \bar{\rho}_1 = \alpha_1.$$

(2) For $i = 1, 2, 3, \dots$ repeat steps 3-6.

(3) (Continue the bidiagonalization.)

$$(a) \beta_{i+1} u_{i+1} = A v_i - \alpha_i u_i$$

$$(b) \alpha_{i+1} v_{i+1} = A^T u_{i+1} - \beta_{i+1} v_i.$$

(4) (Construct and apply next orthogonal transformation.)

$$(a) \rho_i = (\bar{\rho}_i^2 + \beta_{i+1}^2)^{1/2}$$

$$(b) c_i = \bar{\rho}_i / \rho_i$$

$$(c) s_i = \beta_{i+1} / \rho_i$$

$$(d) \theta_{i+1} = s_i \alpha_{i+1}$$

$$(e) \bar{\rho}_{i+1} = -c_i \alpha_{i+1}$$

$$(f) \phi_i = c_i \bar{\phi}_i$$

$$(g) \bar{\phi}_{i+1} = s_i \bar{\phi}_i.$$

ACM Transactions on Mathematical Software, Vol 8, No 1, March 1982

(5) (Update x , w .)

$$(a) x_i = x_{i-1} + (\phi_i / \rho_i) w_i$$

$$(b) w_{i+1} = v_{i+1} - (\theta_{i+1} / \rho_i) w_i.$$

(6) (Test for convergence.)

Pseudo-Code

```
//INIT
beta = norm(b);
u = b / beta;
v = At*u;
alpha = norm(v);
v = v/alpha;
w = v;
x = 0;
phi_hat = beta;
rho_hat = alpha;
while(not converged){
    //bidiagonalization
    u = A * v - alpha * u;
    beta = norm(u);
    u = u / beta;
    v = At * u - beta * v;
    alpha = norm(v);
    v = v / alpha;
    // orthogonal transformation
    rho = sqrt(rho_hat^2 + beta^2);
    c = rho_hat / rho;
    s = beta / rho;
    theta = s * alpha;
    rho_hat = -c * alpha;
    phi = c * phi_hat;
    phi_hat = s * phi_hat;
    //update next vectors
    x = x + (phi / rho) * w;
    w = v - (theta / rho) * w;
    residual = norm(A*x - b);
    checkForConvergence();
}
```

Implementation in cuBLAS

```
//Ax - b (result in tempVector)
tempDouble = -1.0;
tempDouble2 = 1.0;
status = cublasDgemv (handle, CUBLAS_OP_N,
cuBLASCheck(status,__LINE__);
status = cublasDnrm2(handle, tempVector,
cuBLASCheck(status,__LINE__);
improvement = prev_err-curr_err;
printf("line: %d size of error: %.10f\n",
if(improvement<eps) counter++; else
if(counter>1000) break;
prev_err = curr_err;
```

Why cuBLAS?

- Wide range of linear algebra functions (All the required Matrix & Vector operations)
- Tested and trusted library from NVIDIA
- Most of the functions have a cuBLAS equivalent (e.g. `cublasDaxpy` -> `cusparseDaxpyi`)
- Allows debugging while still working in dense format
- Requires no special data type or data structure to work with it
- Has both float and double precision function

Testing

tvercaut / LSQR-cpp

<> Code

! Issues 3

🔗 Pull requests

🎬 Actions

📁 Projects

📖 Wiki

🛡 Security

📈 Insights

🔗 master ▾

🔗 1 branch

🔖 0 tags

Go to file

Add file ▾

📄 Code ▾



tvercaut fixed ordered comparison between pointer and zero

a48ab3e on 17 Feb 2020

🕒 19 commits

📁 Source	fixed ordered comparison between pointer and zero	11 months ago
📁 Testing	introduce GetStoppingReasonMessage() to ease printouts	5 years ago
📁 Utilities/Scripts	Manual copy of the files in the NAMIC Sandbox from Luis Ibanez http:...	5 years ago
📄 .gitignore	Initial commit	5 years ago
📄 CMakeLists.txt	Manual copy of the files in the NAMIC Sandbox from Luis Ibanez http:...	5 years ago

Testing

- We are checking:

- The results
- Number of iteration
- Stopping reason
- Norm of final value of residuals
- Norm of final solution

```
std::cout << "Stopped because " << solver.GetStoppingReason() << ": " << solver.GetStoppingReasonMessage() << std::endl;
std::cout << "Used " << solver.GetNumberOfIterationsPerformed() << " Iterations" << std::endl;
std::cout << "Estimate of final value of norm of residuals = " << solver.GetFinalEstimateOfNormOfResiduals() << std::endl;
std::cout << "Estimate of norm of final solution = " << solver.GetFinalEstimateOfNormOfX() << std::endl;
```


- How do we test?

- Predefined matrices and vectors in multiple sizes


Problems

- Different format that are used in different CUDA libraries

Row-major order



Column-major order



1.1. Data layout

For maximum compatibility with existing Fortran environments, the cuBLAS library uses column-major storage, and 1-based indexing. Since C and C++ use row-major storage, applications written in these languages can not use the native array semantics for two-dimensional arrays. Instead, macros or inline functions should be defined to implement matrices on top of one-dimensional arrays. For Fortran code ported to C in mechanical fashion, one may choose to retain 1-based indexing to avoid the need to transform loops. In this case, the array index of a matrix element in row “i” and column “j” can be computed via the following macro

```
#define IDX2F(i,j,ld) (((j)-1)*(ld))+((i)-1)
```

Problems

- Stopping conditions
- Float / double
- Linking problem
- It works for small matrix like 3×3 but not for a matrix 1000×1000
- Converging problems

Our next steps

- Write the implementation with the cuSPARSE library
- Build a parser for sparse matrices, to read from and write to files.
- Increase parallelization and memory usage of the algorithm to improve performance
- Test and verify our results



Do you have any questions?