

Statistical Machine Learning - Exercise 4



TECHNISCHE
UNIVERSITÄT
DARMSTADT


Dominik Marino - 2468378, Pertamina Kunz - 2380210
August 8, 2020

Contents

1 Neural Network	2
1.1 Multi-layer Perceptron	2
1.2 Deep learning	8
2 Support Vector Machines	9
2.1 Definition	9
2.2 Quadratic Programming	9
2.3 Slack Variables	10
2.4 Slack Variables - Solving Dual optimization problem	11
2.5 The Dual Problem	13
2.6 Kernel Trick	13
2.7 Implementation	14
3 Gaussian Processes	20
3.1 GP Regression	20

1 Neural Network

Given the mnist dataset that contains $N_{tr} = 6006$ training and $N_{te} = 1004$ test data points. Each input has a dimension of $784 = 28 \times 28$, corresponding to a flattened image dataset. One example of such dataset is shown below.



```
1 import numpy as np
2 import matplotlib.pyplot as plt
3
4 # Training Dataset
5 Xtrain = np.genfromtxt("mnist_small_train_in.txt", delimiter=',', dtype=int)
6 ytrain = np.genfromtxt("mnist_small_train_out.txt", dtype=int)
7
8 # Test Dataset
9 Xtest = np.genfromtxt("mnist_small_test_in.txt", delimiter=',', dtype=int)
10 ytest = np.genfromtxt("mnist_small_test_out.txt", dtype=int)
11
12 # Check 1 training sample
13 plt.imshow(Xtrain[0,:].reshape((28,28)))
14 plt.title('Training digit:' + str(ytrain[0]))
15 plt.axis("off")
16 plt.show
```

Listing 1: Reading data and displaying one training sample.

1.1 Multi-layer Perceptron

We started by exploring this selected architecture: 2-layer neural network (1 hidden layer), with 300 nodes in the hidden layer, because it is the simplest network that should give the accuracy of at least 95%/ error of at most 5% (on the full dataset)¹.

Formulation

With 1 hidden layer (2-layer NN),

$$y_k(\mathbf{x}) = f^{(2)} \left(\underbrace{W_{k,0}^{(2)} + \sum_{i=1}^h W_{k,i}^{(2)} f^{(1)} \left(\underbrace{\sum_{j=0}^d W_{i,j}^{(1)} x_j}_{z^{(1)}} \right)}_{a^{(1)}} \right)_{z^{(2)}} \quad (1)$$

where $f^{(1)}(\cdot)$ and $f^{(2)}(\cdot)$ are activation functions in the hidden and output layers, respectively, and, in this case,

- $k = 0, \dots, 9$
- $i = 1, \dots, 300$ or $h = 300$
- $j = 0, \dots, 784$ or $d = 784$

¹<http://yann.lecun.com/exdb/mnist/>




Figure 1: Selected NN architecture for the MNIST dataset.

- $x_0 = 1$ and $W_{i,0}^{(1)}, W_{k,0}^{(2)}, \forall i, k$ correspond to the bias.

Activation functions

We tried using **tanh** and **sigmoid** functions for the hidden layer and **softmax** for the output layer:

$$f^{(1)}(z) = \sigma(z) = \frac{1}{1 + e^{-z}}$$

$$\text{or } f^{(1)}(z) = \tanh(z) = \frac{2}{1 + e^{-2z}} - 1$$

$$f_k^{(2)}(\mathbf{z}) = \hat{y}_k = \frac{e^{z_k}}{\sum_{l=0}^9 e^{z_l}}.$$

The derivative of the sigmoid function:

$$\sigma'(z) = \frac{d}{dz} \left(\frac{1}{1 + e^{-z}} \right) = -\frac{1}{(1 + e^{-z})^2} \cdot \frac{d}{dz} (1 + e^{-z}) = \frac{e^{-z}}{(1 + e^{-z})^2} = \dots = \sigma(z) (1 - \sigma(z)). \quad (2)$$

The derivative of the tanh function:

$$f'^{(1)}(z) = \frac{d}{dz} \tanh(z) = 1 - \tanh^2(z) \quad (3)$$

Loss function

Given the true label y_k and using 1-hot encoder,

$$y_k = \begin{cases} 1 & \text{if } y = k \\ 0 & \text{otherwise.} \end{cases} \quad (4)$$

For classification, we use cross-entropy loss function:

Cross entropy loss function

$$L = - \sum_{k=0}^9 y_k \ln(\hat{y}_k) \quad (5)$$

and since only one y_k has the value of 1, the derivative of cross entropy loss function with softmax activation function:

$$L = - \ln(\hat{y}_k) = - \ln \left(\frac{e^{z_k}}{\sum_{l=0}^9 e^{z_l}} \right) = -z_k + \ln \sum_{l=0}^9 e^{z_l}. \quad (6)$$

The derivative for each z_k :

$$\begin{aligned}
 \nabla_{z_i} L &= \nabla_{z_i} \left(-z_k + \ln \sum_{l=0}^9 e^{z_l} \right) = \nabla_{z_i} \ln \sum_{l=0}^9 e^{z_l} - \nabla_{z_i} z_k \\
 &= \frac{1}{\sum_{l=0}^9 e^{z_l}} \nabla_{z_i} \sum_{l=0}^9 e^{z_l} - \nabla_{z_i} z_k \\
 &= \frac{z_i}{\sum_{l=0}^9 e^{z_l}} - \nabla_{z_i} z_k \\
 &= \hat{y}_i - \delta(i) \\
 &= \hat{y}_i - y_i.
 \end{aligned} \tag{7}$$

Forward- propagation

- $\mathbf{a}^{(0)} = \mathbf{x}$ (already includes bias: $\mathbf{x} = [1, x_1, \dots, x_{784}]^\top$)
- $\mathbf{z}^{(1)} = \mathbf{W}^{(1)} \mathbf{x}$
 - $\mathbf{z}^{(1)} = [z_1, \dots, z_{300}]^\top$, where $z_i = \mathbf{w}_i^\top \mathbf{x}$ and $\mathbf{w}_i = [W_{i,0}^{(1)}, W_{i,1}^{(1)}, \dots, W_{i,784}^{(1)}]$, for $i = 1, \dots, 300$.
 - Matrix–vector form:

$$\mathbf{z}^{(1)} = \begin{bmatrix} z_1^{(1)} \\ z_2^{(1)} \\ \vdots \\ z_{300}^{(1)} \end{bmatrix} = \begin{bmatrix} W_{1,0}^{(1)} & W_{1,1}^{(1)} & \cdots & W_{1,784}^{(1)} \\ W_{2,0}^{(1)} & W_{2,1}^{(1)} & \cdots & W_{2,784}^{(1)} \\ \vdots & \vdots & \ddots & \vdots \\ W_{300,0}^{(1)} & W_{300,1}^{(1)} & \cdots & W_{300,784}^{(1)} \end{bmatrix} \begin{bmatrix} 1 \\ x_1 \\ \vdots \\ x_{784} \end{bmatrix}.$$

- $\mathbf{a}^{(1)} = f^{(1)}(\mathbf{z}^{(1)}) = \sigma(\mathbf{z}^{(1)})$
- $\mathbf{z}^{(2)} = \mathbf{b} + \mathbf{W}^{(2)} \mathbf{a}^{(1)}$
- Matrix–vector form:

$$\mathbf{z}^{(2)} = \begin{bmatrix} z_0^{(2)} \\ z_1^{(2)} \\ \vdots \\ z_9^{(2)} \end{bmatrix} = \begin{bmatrix} W_{0,0}^{(2)} \\ W_{1,0}^{(2)} \\ \vdots \\ W_{9,0}^{(2)} \end{bmatrix} + \begin{bmatrix} W_{0,1}^{(2)} & \cdots & W_{0,300}^{(2)} \\ W_{1,1}^{(2)} & \cdots & W_{1,300}^{(2)} \\ \vdots & \ddots & \vdots \\ W_{9,1}^{(2)} & \cdots & W_{9,300}^{(2)} \end{bmatrix} \begin{bmatrix} a_1 \\ \vdots \\ a_{300} \end{bmatrix}.$$

- $\hat{y}_k = f_k^{(2)}(\mathbf{z}^{(2)}) = \frac{e^{z_k}}{\sum_{l=0}^9 e^{z_l}}$, for $k = 0, \dots, 9$.

Backward-propagation

The update rule: $\mathbf{W}^{(\ell)} = \mathbf{W}^{(\ell)} - \alpha \frac{\partial L}{\partial \mathbf{W}^{(\ell)}}$, where α is the learning rate.

- $\frac{\partial L}{\partial \mathbf{W}^{(2)}} = \frac{\partial L}{\partial \mathbf{z}^{(2)}} \frac{\partial \mathbf{z}^{(2)}}{\partial \mathbf{W}^{(2)}} = \frac{\partial L}{\partial \mathbf{a}^{(1)}} \frac{\partial \mathbf{a}^{(1)}}{\partial \mathbf{z}^{(2)}} \frac{\partial \mathbf{z}^{(2)}}{\partial \mathbf{W}^{(2)}}$
- $\frac{\partial L}{\partial \mathbf{b}} = \frac{\partial L}{\partial \mathbf{z}^{(2)}} \frac{\partial \mathbf{z}^{(2)}}{\partial \mathbf{b}} = \frac{\partial L}{\partial \mathbf{z}^{(2)}}$
- $\frac{\partial L}{\partial \mathbf{W}^{(1)}} = \frac{\partial L}{\partial \mathbf{z}^{(1)}} \frac{\partial \mathbf{z}^{(1)}}{\partial \mathbf{W}^{(1)}} = \frac{\partial L}{\partial \mathbf{a}^{(1)}} \frac{\partial \mathbf{a}^{(1)}}{\partial \mathbf{z}^{(1)}} \frac{\partial \mathbf{z}^{(1)}}{\partial \mathbf{W}^{(1)}} = \frac{\partial L}{\partial \mathbf{z}^{(2)}} \frac{\partial \mathbf{z}^{(2)}}{\partial \mathbf{a}^{(1)}} \frac{\partial \mathbf{a}^{(1)}}{\partial \mathbf{z}^{(1)}} \frac{\partial \mathbf{z}^{(1)}}{\partial \mathbf{W}^{(1)}}$
 - From (7): $\frac{\partial L}{\partial \mathbf{z}^{(2)}} = \hat{\mathbf{y}} - \mathbf{y}|_{10 \times 1}$
 - $\frac{\partial \mathbf{z}^{(2)}}{\partial \mathbf{W}^{(2)}} = \frac{\partial}{\partial \mathbf{W}^{(2)}} (\mathbf{b} + \mathbf{W}^{(2)} \mathbf{a}^{(1)}) = \mathbf{a}^{(1)\top} \Big|_{1 \times 300}$
 - $\frac{\partial \mathbf{z}^{(2)}}{\partial \mathbf{a}^{(1)}} = \frac{\partial}{\partial \mathbf{a}^{(1)}} (\mathbf{b} + \mathbf{W}^{(2)} \mathbf{a}^{(1)}) = \mathbf{W}^{(2)} \Big|_{300 \times 10}$
 - From (2) and (3): $\frac{\partial \mathbf{a}^{(1)}}{\partial \mathbf{z}^{(1)}} = \sigma'(z^{(1)})$ or $\tanh'(z^{(1)}) \Big|_{300 \times 1}$
 - $\frac{\partial \mathbf{z}^{(1)}}{\partial \mathbf{W}^{(1)}} = \mathbf{x}^\top \Big|_{1 \times 785}$.

Then the gradients

- $\left. \frac{\partial L}{\partial \mathbf{W}^{(2)}} \right|_{10 \times 300} = \frac{\partial L}{\partial \mathbf{z}^{(2)}} \frac{\partial \mathbf{z}^{(2)}}{\partial \mathbf{W}^{(2)}} = \hat{\mathbf{y}} - \mathbf{y}|_{10 \times 1} \cdot \mathbf{a}^{(1)\top} \Big|_{1 \times 300}$

- With dimension corrections,

$$\begin{aligned} \left. \frac{\partial L}{\partial \mathbf{W}^{(1)}} \right|_{300 \times 785} &= \frac{\partial \mathbf{z}^{(2)}}{\partial \mathbf{a}^{(1)}} \frac{\partial L}{\partial \mathbf{z}^{(2)}} \otimes \frac{\partial \mathbf{a}^{(1)}}{\partial \mathbf{z}^{(1)}} \cdot \frac{\partial \mathbf{z}^{(1)}}{\partial \mathbf{W}^{(1)}} \\ &= \left(\underbrace{\mathbf{W}^{(2)\top}}_{300 \times 10} \underbrace{(\hat{\mathbf{y}} - \mathbf{y})}_{10 \times 1} \right) \otimes \underbrace{\frac{\partial \mathbf{a}^{(1)}}{\partial \mathbf{z}^{(1)}}}_{300 \times 1} \cdot \underbrace{\mathbf{x}^\top}_{1 \times 785}, \end{aligned}$$

where \otimes is element-wise multiplication.

Preprocessing

$x_0 = 1$ is added to each sample in the training and test set. The y -values are given in all possible numbers $0, 1, \dots, 9$. We first convert this into one-hot-encoding, such that

$$0 \rightarrow [1, 0, 0, 0, 0, 0, 0, 0, 0, 0] \quad \dots \quad 9 \rightarrow [0, 0, 0, 0, 0, 0, 0, 0, 0, 1]$$

Initialization

For the weight matrices, Xavier's initialization is used.²

```

18 ## PREPROCESSING
19 Ntr = ytrain.shape[0] # No. of training samples
20 Nte = ytest.shape[0] # No. of training samples
21 K = 10 # No. of classes
22 L1 = 785 # No. of nodes in input layer
23 L2 = 300 # No. of nodes in hidden layer
24
25 # Add x0=1
26 Xtrain = np.hstack((np.ones((Ntr,1)), Xtrain))
27 Xtest = np.hstack((np.ones((Nte,1)), Xtest))
28 # One-Hot-Encoding
29 Ytrain = np.eye(K)[ytrain]
30 Ytest = np.eye(K)[ytest]
31
32 ## Initialization
33 np.random.seed(1)
34 # weight matrices and bias
35 W1 = np.random.rand(L2,L1) / np.sqrt(L1)
36 W2 = np.random.rand(K,L2) / np.sqrt(L2)
37 bias = np.random.rand(K,1) / np.sqrt(K)

```

Listing 2: Preprocessing and initialization.

Training and Testing

The weight updates are done in mini batches (averaged per mini batch). For each epoch run, the training samples are first randomized. The different learning rates as seen in the results in Figure 2 were chosen simply because they seem to work, we tried different values but they failed to improve the training errors after many epochs/very slow.

With minibatch size of 256 and learning rates of 1 and 0.1 respectively for first->hidden and hidden->output layers, we obtained the best performance after about 1000 epochs: testing error of $66/1004 \times 100\% \approx 6.6\%$. Beyond 1000 epochs, the network started to be overfitting to the training data, as seen in Figure 3, where training error keeps declining while the testing error starts increasing.

With only one fully-connected (dense) hidden layer, it takes many epochs to train the network to get $<8\%$ test error. This should be faster if convolutional NN or more hidden layers were to be used.

```

39 #####
40 def tanh_grad(a):
41     return 1-(np.tanh(a))**2
42 def sigmoid(a):
43     return np.where(a>= 0, 1/(1 + np.exp(-a)), np.exp(a)/(1 + np.exp(a)))
44 def sig_grad(a):
45     return sigmoid(a) * (1 - sigmoid(a))
46 def softmax(a):
47     e_a = np.exp(a - np.max(a)) # max(a) for numerical stability
48     return e_a / e_a.sum(axis=0)
49 def Loss(y, y_hat):
50     return np.sum(y * np.log(y_hat + 1e-6)) # add 1e-6 for numerical stability
51
52 ## TRAINING & TESTING
53 ## ==> WARNING!! SLOW!! NOT OPTIMIZED WHATSOEVER, Works though!
54 epochs = 20000 # How many complete passes through the training set

```

²<https://towardsdatascience.com/weight-initialization-techniques-in-neural-networks-26c649eb3b78>




Figure 2: With learning rates of 1 and 0.1 for first->hidden and hidden->output layers, respectively




Figure 3: Closer look: overfitting beyond 1000 epochs.

```

55
56 batchsize = 256 # just because
57 alpha1 = 1 # Learning rate
58 alpha2 = 0.1 # Learning rate
59 dL_dW1 = 0
60 dL_dW2 = 0
61 i=0
62 train_loss = []
63 TR = []
64 TE = []
65 while i<epochs:
66     idx = np.random.permutation(Ntr)
67     X = Xtrain[idx,:]
68     Y = Ytrain[idx,:]
69
70     # Train with mini batches
71     for start_idx in range(0, Ntr, batchsize):
72         end_idx = min(start_idx+batchsize, Ntr)
73
74         # mini batch
75         Xb = np.transpose(X[start_idx:end_idx,:])
76         Yb = np.transpose(Y[start_idx:end_idx,:])
77
78         minibatch_grads1 = np.empty(W1.shape)
79         minibatch_grads2 = np.empty(W2.shape)
80         minibatch_gradsb = np.empty((K,1))
81         minibatch_loss = np.empty(Xb.shape[1])
82         for b in np.arange(Xb.shape[1]):
83             # Forward feed
84             # input => hidden layer
85             Z1 = np.dot(W1, Xb[:,b])[0, np.newaxis]
86             #a1 = sigmoid(Z1)
87             a1 = np.tanh(Z1)
88             # hidden layer => output
89             Z2 = bias + np.matmul(W2,a1)
90             Y_hat = softmax(Z2)
91
92             # Backward feed/ Gradients
93             dL_dz2 = (Y_hat - Yb[:,b])[0, np.newaxis]
94             #da1_dz1 = sig_grad(Z1)

```

```

95     da1_dz1 = tanh_grad(Z1)
96     dL_dW2 = np.matmul( dL_dz2 , a1.T )
97     dL_dW1 = np.matmul( (np.matmul(W2.T, dL_dz2) * da1_dz1), np.transpose(Xb[:,b][:, np.newaxis]))
98
99     minibatch_grads1 += dL_dW1
100    minibatch_grads2 += dL_dW2
101    minibatch_gradsb += dL_dz2
102    #minibatch_loss[b] = Loss(Yb[:,b], Y_hat)
103
104    # Update
105    dW1 = minibatch_grads1 / Xb.shape[1]
106    dW2 = minibatch_grads2 / Xb.shape[1]
107    db = minibatch_gradsb / Xb.shape[1]
108    #train_loss.append(np.mean(minibatch_loss))
109    W1 = W1 - alpha1 * dW1
110    W2 = W2 - alpha2 * dW2
111    bias = bias - alpha2 * db
112
113    # Training error
114    tr_e = 0
115    for n in range(Ntr):
116        x = Xtrain[n,:].T
117        Z1 = np.dot(W1, x)[:, np.newaxis]
118        #a1 = sigmoid(Z1)
119        a1 = np.tanh(Z1)
120        Z2 = bias + np.matmul(W2,a1)
121        Y_hat = softmax(Z2)
122        if np.argmax(Y_hat)!=ytrain[n]:
123            tr_e += 1
124
125    # Test error
126    te_e = 0
127    for n in range(Nte):
128        x = Xtest[n,:].T
129        Z1 = np.dot(W1, x)[:, np.newaxis]
130        #a1 = sigmoid(Z1)
131        a1 = np.tanh(Z1)
132        Z2 = bias + np.matmul(W2,a1)
133        Y_hat = softmax(Z2)
134        if np.argmax(Y_hat)!=ytest[n]:
135            te_e += 1
136
137    print("epoch #"+str(i+1)+" Tr err: "+str(tr_e) +", Te err: "+str(te_e))
138    TR.append(tr_e)
139    TE.append(te_e)
140    i+=1

```

Listing 3: Training and Testing.

```

142 TR_perc = np.array(TR)/Ntr*100
143 TE_perc = np.array(TE)/Nte*100
144 plt.figure(figsize=(15,4))
145 plt.plot(np.arange(1,TR_perc.shape[0]+1), TR_perc, label="Training error percentage")
146 plt.plot(np.arange(1,TE_perc.shape[0]+1), TE_perc, label="Test error percentage")
147 plt.yticks(np.arange(0,90,10))
148 plt.xlabel("Epochs #")
149 plt.ylabel("%")
150 plt.legend()
151 plt.grid('on')
152 plt.show()
153
154 plt.figure(figsize=(15,4))
155 plt.plot(np.arange(601,TR_perc.shape[0]+1), TR_perc[600:], label="Training error percentage")
156 plt.plot(np.arange(601,TE_perc.shape[0]+1), TE_perc[600:], label="Test error percentage")
157 plt.yticks(np.arange(0,10,1))
158 plt.xlabel("Epochs #")
159 plt.ylabel("%")
160 plt.legend()
161 plt.title("Closer look")
162 plt.grid('on')
163 plt.show()

```

Listing 4: Plotting the percentages.

1.2 Deep learning

- **Highlight the qualitative differences between classical neural networks and deep networks.**

While both classical/shallow neural networks and deep networks try to imitate how human nerve cells work, deep networks usually involve many more hidden layers to uncover some patterns/relationships in the features. The architecture of deep networks is more complex, but the concepts stay the same. In a deep network, the function value y is a composition of many-level function

$$y = (f_K \circ f_{K-1} \circ \dots \circ f_1)(x),$$

where K is typically >3 , x are the inputs (e.g. images), and y are the observations (e.g. class labels). By adding more layers and nodes, a deep network can represent more complex functions and learn more features compared to a shallow network.

- **Which limitations of classical NN does deep learning overcome? Give an intuition of the innovations introduced in deep learning compared to traditional NN.**

In classical NN, one may need to perform labor-intensive feature-engineering. In deeper NN, the network depends less on feature engineering; it will extract and organize the discriminative information from the data. Deep learning methods showed promise in using the data itself to learn new good features and have been very successful in signal processing, computer vision, speech recognition, and natural language processing.

Deep networks promote the re-use of features, and can potentially lead to more abstract features. Deep learning enables representation learning, which can be convenient to express general priors (smoothness, sparsity, temporal and spatial coherence, etc) about the world which would be likely to be useful to solve AI-tasks.

In deeper networks, one can forego using fully connected (dense) layers. Dense layer may be prohibitive in terms of computation time without improving the accuracy of the model. In deep networks, one can add convolutional/recurrent networks in the layers closer to the input which normally have much a lot more nodes. The fewer connections help improving the performance. This is made possible through convolutional and recurrent networks for taking into account spatial and temporal relationships among the inputs. With deeper networks it is also possible to incorporate regularization techniques like Dropout to randomly "kill" some nodes to avoid overfitting.

- **Are deeper networks necessarily better from a theoretical perspective?**

In a shallow network you may need to train the weights a lot more times to reach the accuracy (or even lower accuracy) of a deeper network. However, a deeper network also has high computational cost due to a lot more connections which means a lot of weight to update. Due to the millions of parameters, getting enough labeled data to train is difficult. Furthermore, a deeper network may fit to the data on hand too well, i.e. overfitting. Therefore, it is not necessarily better.

3 4 5 6

³Goodfellow, I., Bengio, Y. and Courville, A., 2016. Deep learning. MIT press.

⁴Murphy, Kevin P. Machine learning: a probabilistic perspective. MIT press, 2012.

⁵Bengio, Yoshua, Aaron Courville, and Pascal Vincent. "Representation learning: A review and new perspectives." IEEE transactions on pattern analysis and machine intelligence 35.8 (2013): 1798-1828.

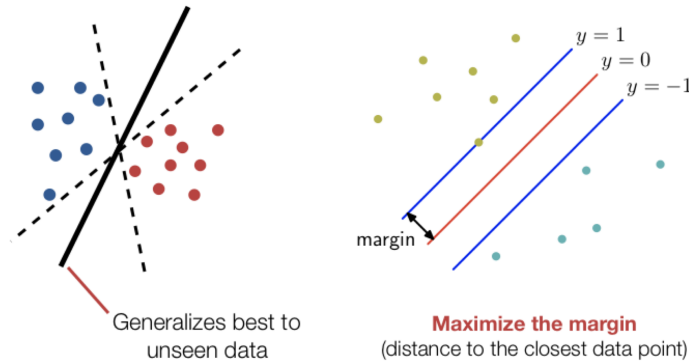
⁶Deisenroth, A. Aldo Faisal Marc Peter, and Cheng Soon Ong. "Mathematics for Machine Learning. 2019." url: <https://mml-book.github.io/> io: 407.

2 Support Vector Machines

2.1 Definition

- Definition of SVM

- Support Vector Machine is a supervised machine learning algorithm which can classification and regression challenges. SVM is a margin classifier and supports linear and non-linear classification. With the SVM algorithm, we can plot each data in a n-dimensional space (n is the number of features) and then perform by finding the best hyperplane do seperate two or more classes. SMV tries to find the best margin (distance between the line and the support vector) that seperates the classes and reduce the risk of an error in the data⁷.



- Advantages with respect to (w.r.t) other linear approches

- SVM is more effective in high dimensional space. Meanwhile in regression the gram matrix $X^T X$ is not invertible. The linear model $\hat{\beta} = (X^T X)^{-1} X^T y$ completely fails. To improve this model, we can use ridge regression with additional penalty.

$$\hat{\beta}^{Ridge} = \operatorname{argmin}_{\beta \in R} \|Y - X\|_2^2 + \lambda \|\beta\|_2^2$$

Logistic regresssion focus on maximinzng the probability of the data and SVM tries to find the seperating hyperplane that maximizes the distance of the closest points to the margin. SVM dont penalize examples for which the correct decision is made with sufficient confidence. This may SVMs be good for generalization and a robust method against outliers.

The biggest advantages of using Support Vector Machine over a linear approach is to classify non-linear datasets with non-linear kernels. SVMs depends on their support vectors and not on the entire dataset for solving the decision boundary.

2.2 Quadratic Programming

Constrained optimization problem:

$$\begin{aligned} \operatorname{argmin}_{\mathbf{w}, b} \quad & \frac{1}{2} \|\mathbf{w}\|^2 \\ \text{s.t.} \quad & y_i(\mathbf{w}^T x_i + b) - 1 \geq 0 \quad \forall i \end{aligned}$$

Langrangian Formulation:


$$L(\mathbf{w}, b, \alpha) = \frac{1}{2} \|\mathbf{w}\|^2 - \sum_{i=1}^N \alpha_i (y_i(\mathbf{w}^T x_i + b) - 1)$$

⁷Picture Source: SML Lecture Script 11 Slide 8

2.3 Slack Variables

- Explain concept of slack variables

- In a linear programming problem is a slack variable the inequality represent the constraint that can be replaced by the equations. The slack variable cannot take negative values. It is always required to be positive or zero. If the slack variable is zero for the data-point, then there are correctly classified and are either on the margin or on the correct side of the margin. If the points $0 < \epsilon < 1$ there lie inside the margin, but on the correct side of the decision boundary. If the data-point for the $\epsilon > 0$ they lie on the wrong side of the decision boundary and they are missclassified. Our goal is to maximize the margin while softly penalizing points that lie on the wrong side of the margin boundary. Shown in the graphic⁸.



$$\frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{i=1}^n \epsilon_i$$

Parameters:

- * C: Costparameter for the missinformation. It is a regularization coefficient because it contrails the trade-off between minimizing training errors and controlling model complexity.
- * ϵ : non-negative slack variable $\rightarrow \sum_n \epsilon_n$: is a upperbound on the number of misclassified points
- * \mathbf{w} : targetfunction
- * b: bias term

- Reformulate the optimizaiton problem

$$\begin{aligned} \operatorname{argmin}_{\mathbf{w}, b, \epsilon} \quad & \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{i=1}^n \epsilon_i \\ \text{s.t.} \quad & y_i(\mathbf{w}^T x_i + b) \geq 1 - \epsilon_i \\ \text{s.t.} \quad & \epsilon_i \geq 0 \end{aligned}$$

Langrangian Formulation:

$$\begin{aligned} L(\mathbf{w}, b, \alpha, \beta) = & \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{i=1}^n \epsilon_i - \sum_{i=1}^N \alpha_i (y_i(\mathbf{w}^T x_i + b) - 1 + \epsilon_i) - \sum_{i=1}^N \beta_i \epsilon_i \\ \text{with } & \alpha_i > 0, \beta_i > 0 \end{aligned}$$

⁸Picture Source: Pattern Recognition and Machine Learning page 332

2.4 Slack Variables - Solving Dual optimization problem

Langrangian Formulation:

$$L(\mathbf{w}, b, \alpha, \beta) = \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{i=1}^n \epsilon_i - \sum_{i=1}^N \alpha_i (y_i (\mathbf{w}^T x_i + b) - 1 + \epsilon_i) - \sum_{i=1}^N \beta_i \epsilon_i$$

with $\alpha_i > 0, \beta_i > 0$

Set deviate of L w.r.t w and b to zero:

$$\begin{aligned} \nabla_w L \stackrel{!}{=} 0 &\Rightarrow w - \sum_{i=1}^N \alpha_i \beta_i x_i \stackrel{!}{=} 0 \\ &\Rightarrow w = \sum_{i=1}^N \alpha_i \beta_i x_i \end{aligned}$$

$$\begin{aligned} \nabla_b L \stackrel{!}{=} 0 &\Rightarrow - \sum_{i=1}^N (\alpha_i \beta_i w^T x_i + \alpha_i y_i b - \alpha_i + \epsilon_i \alpha_i) \stackrel{!}{=} 0 \\ &\Rightarrow - \sum_{i=1}^N \alpha_i y_i = 0 \\ &\Rightarrow \sum_{i=1}^N \alpha_i y_i = 0 \end{aligned}$$

$$\begin{aligned} \nabla_\epsilon L \stackrel{!}{=} 0 &\Rightarrow C \sum_{i=1}^N 1 - \sum_{i=1}^N \beta_i - \sum_{i=1}^N \alpha_i \\ &\Rightarrow C * N - N * \beta_i - N * \alpha_i = 0 \\ &\Rightarrow \alpha_i = C - \beta_i \end{aligned}$$

$$\nabla_\alpha L \stackrel{!}{=} 0 \Rightarrow \sum_{i=1}^N y_i (w^T x_i + b) - 1 + \epsilon_i = 0$$

$$\nabla_\beta L \stackrel{!}{=} 0 \Rightarrow \sum_{i=1}^N \epsilon_i = 0$$

Substituting in Lagrangian:

$$\begin{aligned} L(\mathbf{w}, b, \alpha, \beta) &= \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{i=1}^N \epsilon_i - \sum_{i=1}^N \alpha_i (y_i (\mathbf{w}^T x_i + b) - 1 + \epsilon_i) - \sum_{i=1}^N \beta_i \epsilon_i \\ &= \frac{1}{2} w^T w + C \sum_{i=1}^N \epsilon_i - \sum_{i=1}^N \alpha_i y_i w^T x_i - \sum_{i=1}^N \alpha_i y_i b + \sum_{i=1}^N \alpha_i - \sum_{i=1}^N \alpha_i \epsilon_i - \sum_{i=1}^N \beta_i \epsilon_i \end{aligned}$$

With the derivation we can simplify the terms:

$$\begin{aligned}
L(\mathbf{w}, b, \alpha, \beta) &= \frac{1}{2}w^T w + C \sum_{i=1}^N \epsilon_i - \sum_{i=1}^N \alpha_i y_i w^T x_i - \sum_{i=1}^N \alpha_i y_i b + \sum_{i=1}^N \alpha_i - \sum_{i=1}^N \alpha_i \epsilon_i - \sum_{i=1}^N \beta_i \epsilon_i \\
&= \sum_{i=1}^N \alpha_i + \frac{1}{2}w^T w - \sum_{i=1}^N \alpha_i y_i w^T x_i - \sum_{i=1}^N \alpha_i \epsilon_i - \sum_{i=1}^N \beta_i \epsilon_i \\
&= \sum_{i=1}^N \alpha_i + w^T \left[\frac{1}{2}w - \sum_{i=1}^N \alpha_i y_i x_i \right] - \sum_{i=1}^N (\alpha_i \epsilon_i + \beta_i \epsilon_i) \\
&= \sum_{i=1}^N \alpha_i + w^T \left[\frac{1}{2}w - \sum_{i=1}^N \alpha_i y_i x_i \right] - \sum_{i=1}^N \epsilon_i (\alpha_i + \beta_i)
\end{aligned}$$

Set $\alpha_i = C - \beta_i$ in:

$$\begin{aligned}
L(\mathbf{w}, b, \alpha, \beta) &= \sum_{i=1}^N \alpha_i + w^T \left[\frac{1}{2}w - \sum_{i=1}^N \alpha_i y_i x_i \right] - \sum_{i=1}^N \epsilon_i (C - \beta_i + \beta_i) \\
&= \sum_{i=1}^N \alpha_i + w^T \left[\frac{1}{2}w - \sum_{i=1}^N \alpha_i y_i x_i \right] - \sum_{i=1}^N \epsilon_i (C) \\
&= \sum_{i=1}^N \alpha_i + w^T \left[\frac{1}{2}w - \sum_{i=1}^N \alpha_i y_i x_i \right] - C \sum_{i=1}^N \epsilon_i \\
&= \sum_{i=1}^N \alpha_i + w^T \left[\frac{1}{2}w - \sum_{i=1}^N \alpha_i y_i x_i \right] - C \sum_{i=1}^N \epsilon_i
\end{aligned}$$

Now we can set the formula for $w = \sum_{i=1}^N \alpha_i \beta_i x_i$ in our Lagrangian from our derivation:

$$\begin{aligned}
L(\mathbf{w}, b, \alpha, \beta) &= \sum_{i=1}^N \alpha_i + w^T \left[\frac{1}{2} \sum_{i=1}^N \alpha_i \beta_i x_i - \sum_{i=1}^N \alpha_i y_i x_i \right] \\
&= \sum_{i=1}^N \alpha_i + w^T \left[-\frac{1}{2} \sum_{i=1}^N \alpha_i \beta_i x_i \right] \\
&= \sum_{i=1}^N \alpha_i + \left(\sum_{i=1}^N \alpha_i \beta_i x_i \right)^T \left[-\frac{1}{2} \sum_{i=1}^N \alpha_i \beta_i x_i \right] \\
&= \sum_{i=1}^N \alpha_i + \sum_{i=1}^N \alpha_i \beta_i x_i^T \left[-\frac{1}{2} \sum_{i=1}^N \alpha_i \beta_i x_i \right] \\
&= \sum_{i=1}^N \alpha_i - \frac{1}{2} \sum_{i,j=0}^N \alpha_i \alpha_j y_i y_j x_i^T x_j
\end{aligned}$$

Parameters:

- Scaler:
 - α : Lagrange multiplier for training example
 - y : -1 or 1
- Vector:
 - x : feature vector of training data

Dual Optimization Problem:

$$\begin{aligned} \max_{\alpha} \quad & W(\alpha) = \sum_{i=1}^N \alpha_i - \frac{1}{2} \sum_{i,j=0}^N \alpha_i \alpha_j y_i y_j x_i^T x_j \\ \text{s.t.} \quad & \alpha_i \geq 0 \quad i = 1, \dots, m \\ & \sum_i^m \alpha_i y_i = 0 \end{aligned}$$

Hence, we can solve dual in lieu of primal problem:

$$\max_{\alpha} \quad W(\alpha) = \sum_{i=1}^N \alpha_i - \frac{1}{2} \sum_{i=1}^N \sum_{j=1}^N \alpha_i \alpha_j y_i y_j \langle x_i, x_j \rangle$$

- $\langle x_i, x_j \rangle$: Inner product of feature vector

If we want to solve the features in higher dimensional space, we have to replace the features x by $\phi(x)$

2.5 The Dual Problem

- Advantages of solving the dual instead of the primal?

Primal:

$$\min_{w,b} \max_{\alpha \geq 0} \quad \frac{1}{2} \|\mathbf{w}\|^2 - \sum_{i=1}^N \alpha_i (y_i (\mathbf{w}^T x_i + b) - 1)$$

Dual:

$$\max_{\alpha \geq 0} \min_{w,b} \quad \frac{1}{2} \|\mathbf{w}\|^2 - \sum_{i=1}^N \alpha_i (y_i (\mathbf{w}^T x_i + b) - 1)$$

- The dual problem allows it to reformulate the model using kernels and maximize the margin classifier that can efficiently applied to the feature space. The most significant benefit from solving it dual is the Kernel Trick (explained in the 2.6). The dual Kernel evaluation requires less kernel evaluation than the primal. This gave the dual a more stable solution in less computation time. The training time of the dual SVM is faster than the primal SVM and usually the number of the support vector are smaller. Another advantage is that the dual problem has many variables as constraints as there are in the primal. So we can reduce the dimensionality of the matrix so much⁹.

2.6 Kernel Trick

- Explain the kernel trick and why is it particularly convenient in SVMs
 - The Lagrangian transform the SVM problem to get a good generalization for our data and so we can drop the slack variable. Lagrangian allows us to take the constrained optimization problem and reformulate it as an unconstrained problem, like:

$$L(\mathbf{w}, b, \alpha) = \frac{1}{2} \|\mathbf{w}\|^2 - \sum_{i=1}^N \alpha_i (y_i (\mathbf{w}^T x_i + b) - 1)$$

The minus sign with the second term is used because we are minimizing with respect to the first term, but maximizing the second. Now we can simplify the result, and making use of our constraint and derive our model:

Dual Lagrangian:

$$W(\alpha) = \sum_{i=1}^N \alpha_i - \frac{1}{2} \sum_{i,j=0}^N \alpha_i \alpha_j y_i y_j x_i^T x_j$$


This objective function is actually more expensive to evaluate than the primal Lagrangian. So it leads to following formula form:

$$W(\alpha) = \sum_{i=1}^N \alpha_i - \frac{1}{2} \sum_i \sum_j \alpha_i \alpha_j y_i y_j k(x_i, x_j)$$

⁹Paper from Shigeo Abe: Is primal better than dual https://www.researchgate.net/publication/262407522_Is_Primal_Better_Than_Dual

where $k(x_i, x_j) = \phi(x_i)^T \phi(x_j)$ is a kernel function. With the kernel trick it frees us from thinking about the features directly. We can classify data that is not linearly separable in the original feature's space. In our feature space we don't have to compute the coordinates of the data in the space, but we can easily compute the inner product of all pairs of data in our feature space, because it is a lot of easier to get inner product in a high dimensional space than the actual points in a high dimensional space. It is computationally cheaper than the explicit computation of the coordinates.

2.7 Implementation



We see in our data, that we can use a linear kernel to solve this task. We also get a more generalization of our data to get a nicer hyperplane between Setosa and Virginica. A polynomial kernel would probably overfitting. For this case a non-linear kernel (polynomial, radial basis function, etc.) wouldn't be useful.

```
1 import numpy as np
2 import matplotlib.pyplot as plt
3 from cvxopt import matrix as cvxopt_matrix
4 from cvxopt import solvers as cvxopt_solvers
5
6 iris = np.genfromtxt('dataSets/iris-pca.txt')
7 X_data, y_data = iris[:, 0:2], iris[:, 2]
8
9 def show_data(X, y):
10     plt.figure()
11     plt.title('Iris Dataset --> red=Setosa, yellow=Virginica')
12     plt.scatter(X[:, 0], X[:, 1], c=y, s=50, cmap='autumn')
13     plt.xlabel('x')
14     plt.ylabel('y')
15     plt.show()
```

Listing 5: Iris Dataset


```
120 def main():
121     def change_labels(labels):
122         Y = []
123         for y in labels:
124             if y == 0:
125                 Y.append(-1)
126             else:
127                 Y.append(1)
128         return np.asarray(Y)
```

```

129 X, y = X_data, change_labels(y_data)
130 show_data(X, y)
131 doSMV(X, y, linear_kernel, 10)
132
133
134 if __name__ == '__main__':
135     main()

```

Listing 6: Main



```

31 def doSMV(X, y, kernel, C = 10):
32     samples, features = X.shape
33     K = np.zeros((samples, samples))
34     for i in range(samples):
35         for j in range(samples):
36             K[i, j] = kernel(X[i], X[j])
37     H = np.outer(y, y) * K
38
39     # Converting into cvxopt format
40     P = cvxopt_matrix(H)
41     q = cvxopt_matrix(-1 * np.ones((samples)))
42     G = cvxopt_matrix(np.vstack((np.diag(np.ones(samples)) * -1, np.identity(samples))))
43     h = cvxopt_matrix(np.hstack((np.zeros(samples), np.ones(samples) * C)))
44     A = cvxopt_matrix(y, (1, samples), 'd')
45     b = cvxopt_matrix(0.0)
46
47     #solve optimization problem for quadratic programming
48     cvxopt_solvers.options['show_progress'] = False
49     sol = cvxopt_solvers.qp(P, q, G, h, A, b)
50     alphas = np.ravel(sol['x'])
51
52     # Selecting the set of indices S corresponding to non zero parameters
53     S = alphas > 1e-5
54
55     support_vectors = X[S]
56     support_alphas = alphas[S]
57     support_vector_y = y[S]
58
59
60     w = ((y * alphas).T @ X).reshape(-1, 1)
61     b = np.mean(support_vector_y - np.dot(support_vectors, w))
62

```

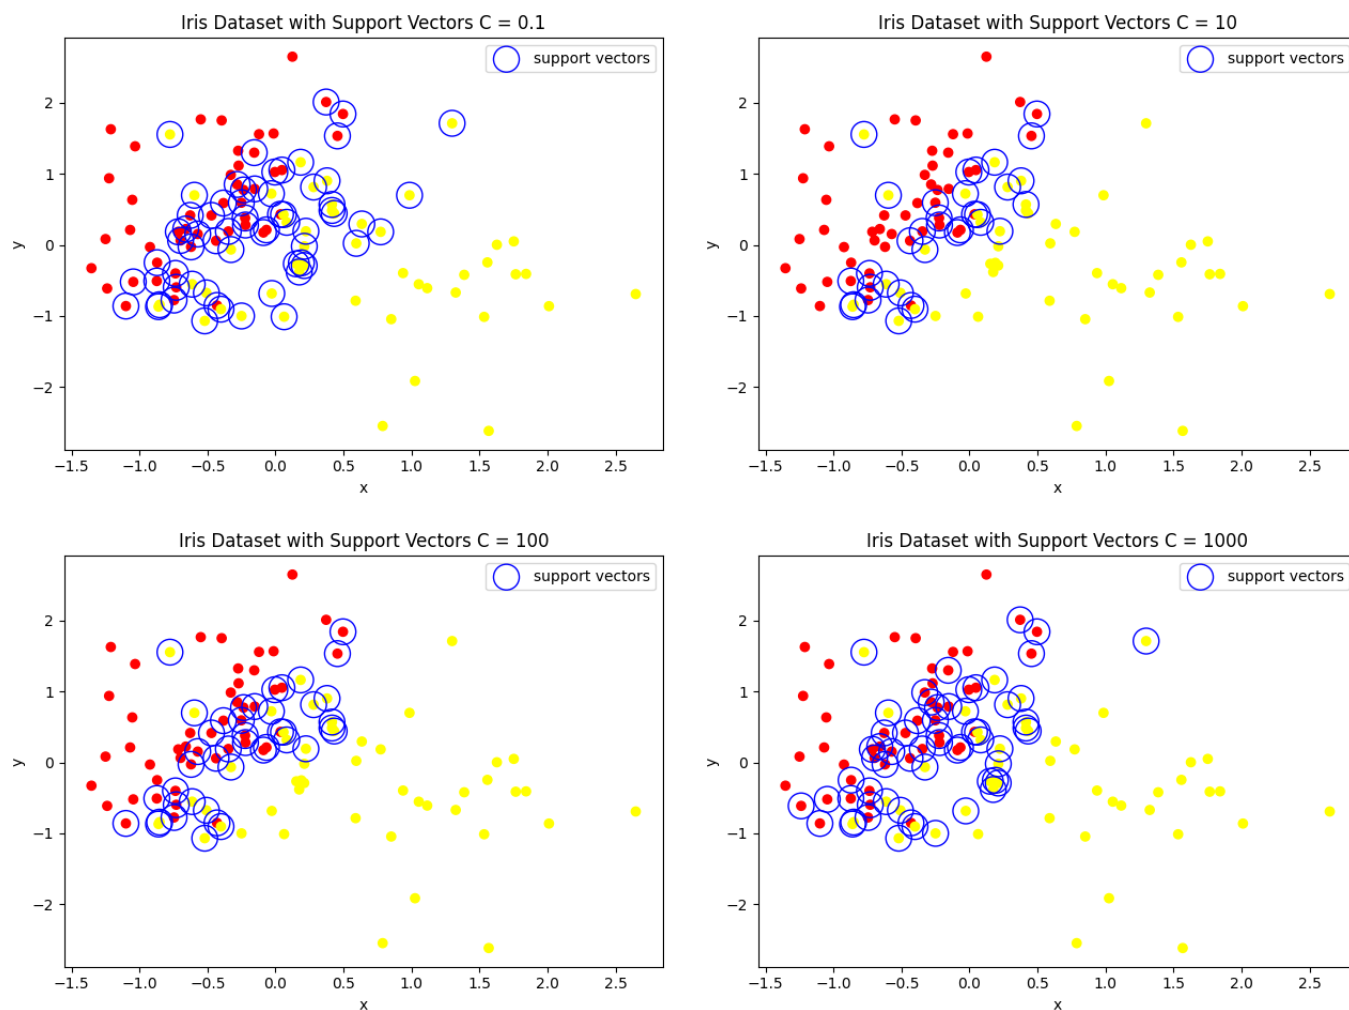
```

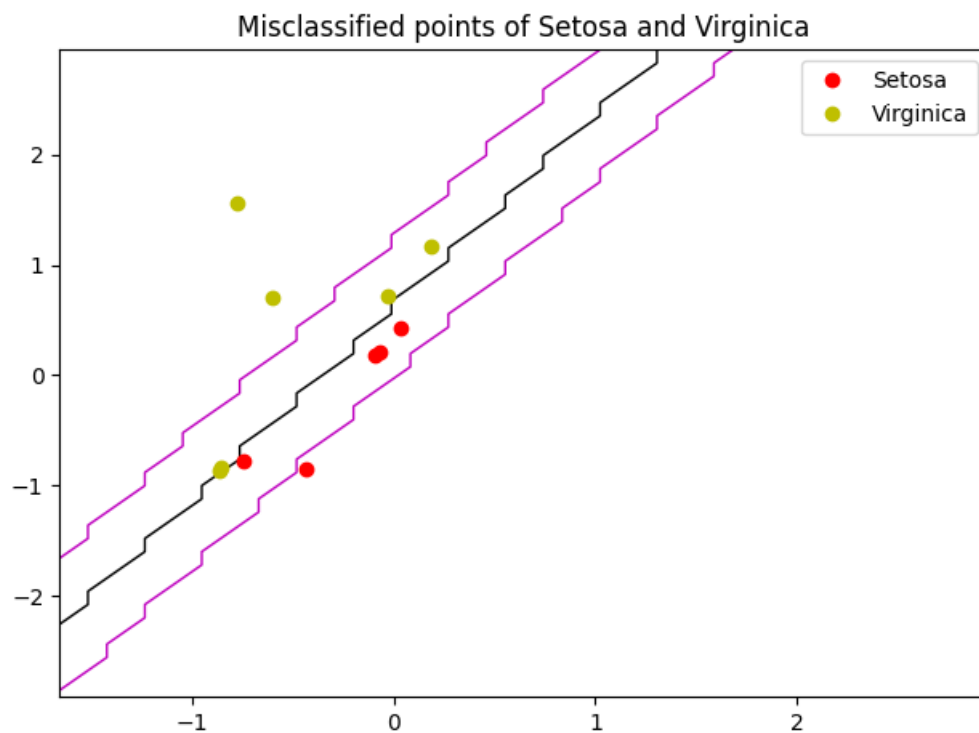
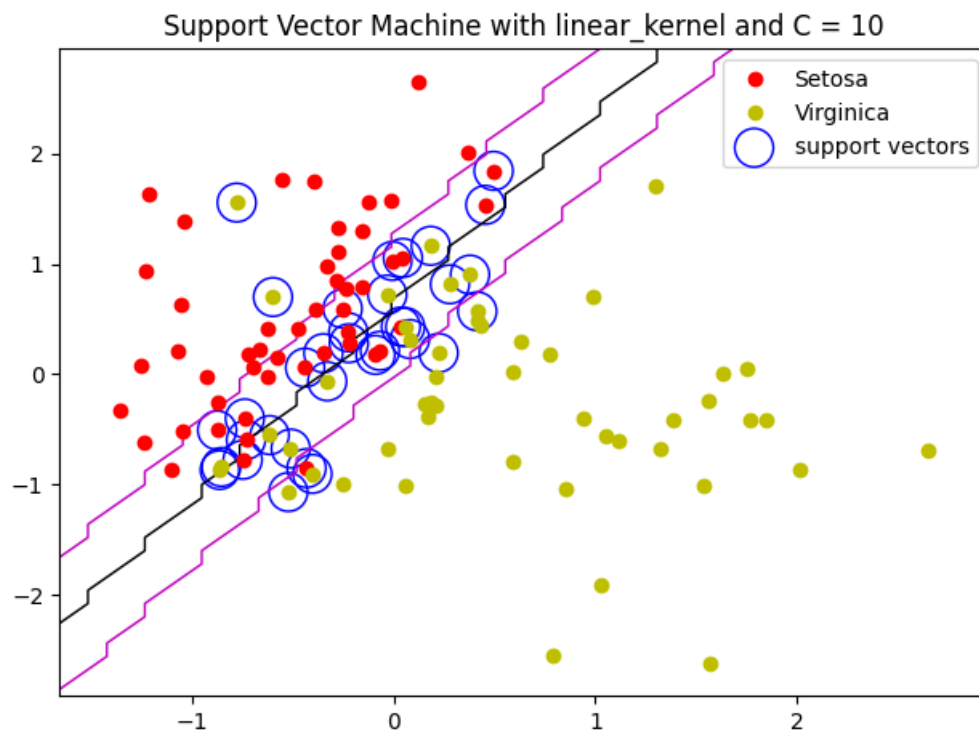
63 #print('Alphas = ', support_alphas)
64 #print('w = ', w.flatten())
65 #print('b = ', b)
66
67 plot_data_with_support_vectors(X, y, support_vectors, C)
68 show_SMV_Plot(w, b, X, y, support_vectors, kernel.__name__, C)
69

```

Listing 7: Optimization of SVM

Evaluating different cost-factors C for the misclassification and get the perfect value that fits the dataset, we have to split the dataset into a Trainingset and Testset and cross-validate it on the data. Calculate the Error and interpret the results (not doing it here, we only say that $C=10$ fits at best):





```

1 import numpy as np
2 import matplotlib.pyplot as plt
3 from cvxopt import matrix as cvxopt_matrix
4 from cvxopt import solvers as cvxopt_solvers
5
6 iris = np.genfromtxt('dataSets/iris-pca.txt')
7 X_data, y_data = iris[:, 0:2], iris[:, 2]

```

```

8
9 def show_data(X, y):
10     plt.figure()
11     plt.title('Iris Dataset --> red=Setosa, yellow=Virginica')
12     plt.scatter(X[:, 0], X[:, 1], c=y, s=50, cmap='autumn')
13     plt.xlabel('x')
14     plt.ylabel('y')
15     plt.show()
16
17 def linear_kernel(x, y):
18     return np.dot(x, y)
19
20 def plot_data_with_support_vectors(X, y, support_vectors, C):
21     plt.figure()
22     plt.title('Iris Dataset with Support Vectors C = {}'.format(C))
23     plt.scatter(X[:, 0], X[:, 1], cmap='autumn', c=y_data)
24     plt.scatter(support_vectors[:, 0], support_vectors[:, 1], label='support vectors', facecolors='none', linewidths=1,
25                 edgecolors='b', s=300)
26     plt.legend()
27     plt.xlabel('x')
28     plt.ylabel('y')
29     plt.show()
30
31 def doSMV(X, y, kernel, C = 10):
32     samples, features = X.shape
33     K = np.zeros((samples, samples))
34     for i in range(samples):
35         for j in range(samples):
36             K[i, j] = kernel(X[i], X[j])
37     H = np.outer(y, y) * K
38
39     # Converting into cvxopt format
40     P = cvxopt_matrix(H)
41     q = cvxopt_matrix(-1 * np.ones((samples)))
42     G = cvxopt_matrix(np.vstack((np.diag(np.ones(samples)) * -1, np.identity(samples))))
43     h = cvxopt_matrix(np.hstack((np.zeros(samples), np.ones(samples) * C)))
44     A = cvxopt_matrix(y, (1, samples), 'd')
45     b = cvxopt_matrix(0.0)
46
47     #solve optimization problem for quadratic programming
48     cvxopt_solvers.options['show_progress'] = False
49     sol = cvxopt_solvers.qp(P, q, G, h, A, b)
50     alphas = np.ravel(sol['x'])
51
52     # Selecting the set of indices S corresponding to non zero parameters
53     S = alphas > 1e-5
54
55     support_vectors = X[S]
56     support_alphas = alphas[S]
57     support_vector_y = y[S]
58
59
60     w = ((y * alphas).T @ X).reshape(-1, 1)
61     b = np.mean(support_vector_y - np.dot(support_vectors, w))
62
63
64     #print('Alphas = ', support_alphas)
65     #print('w = ', w.flatten())
66     #print('b = ', b)
67
68     plot_data_with_support_vectors(X, y, support_vectors, C)
69     show_SMV_Plot(w, b, X, y, support_vectors, kernel.__name__, C)
70
71 def predict(w, X, b, margin = 0):
72     return np.sign(np.dot(X, w) + b + margin)
73
74 def seperate_data(X, labels):
75     return X[labels == -1], X[labels == 1]
76
77 def show_SMV_Plot(w, b, X, y, support_vectors, kernel_name, C):
78     y_pred = predict(w, X, b)
79     correct = np.sum(np.equal(y_pred, y.reshape(-1, 1)))
80     print("{} out of {} predictions are correct".format(correct, len(y_pred)))
81
82     X1, X2 = seperate_data(X, y)
83     plt.plot(X1[:, 0], X1[:, 1], "ro", label='Setosa')
84     plt.plot(X2[:, 0], X2[:, 1], "yo", label='Virginica')
85     plt.scatter(support_vectors[:, 0], support_vectors[:, 1], label='support vectors', facecolors='none', linewidths=1,
86                 edgecolors='b', s=300)

```

```

87 X1, X2 = np.meshgrid(np.linspace(min(X[:, 0]) - .3, max(X[:, 0]) + .3, 50), np.linspace(min(X[:, 1]) - .3, max(X[:,
88 1]) + .3, 50))
89 X = np.array([[x1, x2] for x1, x2 in zip(np.ravel(X1), np.ravel(X2))])
90 Z = predict(w, X, b).reshape(X1.shape)
91 plt.contour(X1, X2, Z, [0.0], colors='k', linewidths=1, origin='lower')
92 plt.contour(X1, X2, predict(w, X, b, 1).reshape(X1.shape), [0.0], colors='m', linewidths=1, origin='lower', )
93 plt.contour(X1, X2, predict(w, X, b, -1).reshape(X1.shape), [0.0], colors='m', linewidths=1, origin='lower')
94 plt.title('Support Vector Machine with ' + kernel_name + ' and C = {}'.format(C))
95 plt.legend()
96 plt.show()
97
98 #misclassified points
99 setosa, virginica = [], []
100 for i, j, k in zip(y, y_pred, X_data):
101     if(i != j):
102         if i == -1:
103             setosa.append(k)
104         else:
105             virginica.append(k)
106 setosa = np.asarray(setosa)
107 virginica = np.asarray(virginica)
108 print(setosa, virginica)
109
110 plt.figure()
111 plt.plot(setosa[:, 0], setosa[:, 1], "ro", label='Setosa')
112 plt.plot(virginica[:, 0], virginica[:, 1], "yo", label='Virginica')
113 plt.contour(X1, X2, Z, [0.0], colors='k', linewidths=1, origin='lower')
114 plt.contour(X1, X2, predict(w, X, b, 1).reshape(X1.shape), [0.0], colors='m', linewidths=1, origin='lower', )
115 plt.contour(X1, X2, predict(w, X, b, -1).reshape(X1.shape), [0.0], colors='m', linewidths=1, origin='lower')
116 plt.legend()
117 plt.title('Misclassified points of Setosa and Virginica')
118 plt.show()
119
120 def main():
121     def change_labels(labels):
122         Y = []
123         for y in labels:
124             if y == 0:
125                 Y.append(-1)
126             else:
127                 Y.append(1)
128         return np.asarray(Y)
129
130     X, y = X_data, change_labels(y_data)
131     show_data(X, y)
132     doSMV(X, y, linear_kernel, 10)
133
134 if __name__ == '__main__':
135     main()

```

Listing 8: Full Code

3 Gaussian Processes

3.1 GP Regression

Target function $y = \sin(x) + \sin^2(x)$ with $x \in [0, 0.005, 0.01, 0.015, \dots, 2\pi]$

The generative model: $t_n = y(x_n) + \epsilon_n, \epsilon \sim \mathcal{N}(0, \sigma_n^2)$.

The prediction for new data points: $p(t_{n+1}) = \mathcal{N}(t_{n+1} | \mathbf{0}, \mathbf{C}_{n+1})$ where

$$\mathbf{C}_{n+1} = \begin{pmatrix} \mathbf{C}_n & \mathbf{k} \\ \mathbf{k} & c \end{pmatrix}$$
$$C_n = k(x_n, x_n) + \sigma_n^2 \delta_{ij}$$
$$\mathbf{k} = [k(x_1, x_{n+1}) \dots k(x_n, x_{n+1})]^\top$$
$$c = k(x_{n+1}, x_{n+1}) + \beta^{-1}$$

The squared exponential kernel

$$k(\mathbf{x}_i, \mathbf{x}_j) = \sigma_f^2 \exp\left(-\frac{1}{2l^2} (x_i - x_j)^2\right)$$

where σ_f^2 is the signal variance, l is the length-scale and σ_n^2 is the noise variance.

The prediction equations for the new data points:

$$m(x_{n+1}) = \mathbf{k}^\top \mathbf{C}_n^{-1} \mathbf{t}$$
$$\sigma^2(x_{n+1}) = c - \mathbf{k}^\top \mathbf{C}_n^{-1} \mathbf{k}$$

Figure 4 shows the prediction results for iterations 1,2,4,8, and 16. Every iteration, the target point with the highest uncertainty is sampled.

```
1 import numpy as np
2 import matplotlib.pyplot as plt
3
4 def kernel(xi, xj, sigma_f=1.0, l=1):
5     return sigma_f * np.exp(-(np.linalg.norm(xi - xj)**2) / (2 * l**2))
6
7 def f(x):
8     return np.sin(x) + np.sin(x)**2
9
10 def Cnew_comps(x, xnew, sigma_f=1.0, l=1): #Calculate the components of Cnew
11     n = x.shape[0]
12     nnew = xnew.shape[0]
13
14     Cn = np.empty((n, n))
15     for i in range(n):
16         for j in range(n):
17             Cn[i,j] = kernel(x[i], x[j], sigma_f=sigma_f, l=1)
18
19     k = np.empty((nnew, n))
20     for i in range(nnew):
21         for j in range(n):
22             k[i,j] = kernel(xnew[i], x[j], sigma_f=sigma_f, l=1)
23
24     c = np.empty((nnew, nnew))
25     for i in range(nnew):
26         for j in range(nnew):
27             c[i,j] = kernel(xnew[i], xnew[j], sigma_f=sigma_f, l=1)
28
29     return (Cn, k, c)
30
31 def gpr_params(Cn, c, k, sigma_n, t): # gaussian regression parameters
32     n = Cn.shape[0]
33     # mean
34     m = np.dot(k, np.dot(np.linalg.inv(Cn + (sigma_n**2)*np.eye(n)), t.reshape([n, 1])))
35     # Covariance.
36     S = c - np.dot(k, np.dot(np.linalg.inv(Cn + (sigma_n**2)*np.eye(n)), k.T))
37     return (m, S)
38
39 np.random.seed(1)
40 # prediction points
41 xnew = np.arange(0,2,0.005) * np.pi
```

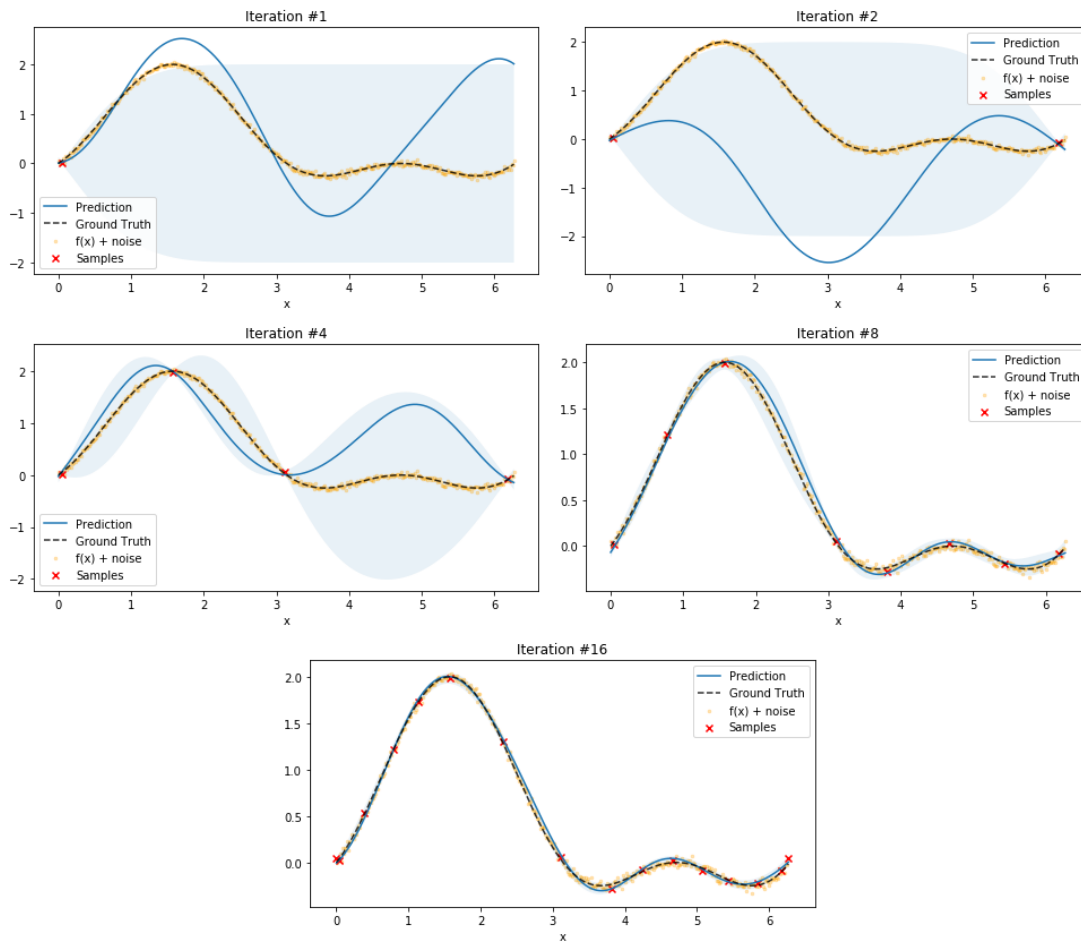


Figure 4: Gaussian process regression, using $l = 1, \sigma_f = 1$. Shaded: $\text{mean} \pm 2 \times \text{standard deviation}$ (uncertainty).

```

42 nnew = xnew.shape[0]
43 # error standard deviation
44 sigma_n = np.sqrt(0.001)
45 # draw error randomly
46 epsilon = np.random.normal(0, sigma_n, nnew)
47 # observed
48 t_observed = f(xnew) + epsilon
49
50 # start with no target data point
51 x = np.array([])
52 t = np.array([])
53 # start with mean=0
54 mnew = np.zeros((1,1))
55 uncertainty = np.random.rand(nnew)
56 for i in range(16):
57     newptidx = np.argmax(uncertainty)
58     x = np.append(x, xnew[newptidx])
59     t = np.append(t, t_observed[newptidx])
60     n = x.shape[0]
61
62 # Cov matrix
63 l = 1
64 sigma_f = 1
65 Cn, k, c = Cnew_comps(x, xnew, sigma_f=sigma_f, l=1)
66 Cnew_left = np.concatenate((Cn + (sigma_n**2)*np.eye(n), k), axis=0)
67 Cnew_right = np.concatenate((k.T, c), axis=0)
68 Cnew = np.concatenate((Cnew_left, Cnew_right), axis=1)
69 # GPR parameters
70 mnew, covnew = gpr_params(Cn, c, k, sigma_n, t)
71 # Sample from multivariate Gaussian with the GPR parameters
72 t_pred = np.random.multivariate_normal(mean=mnew.ravel(), cov=covnew)
73 # Define uncertainty=2stddev

```

```

74     uncertainty = 2 * np.sqrt(np.diag(covnew))
75
76     if i in [0,1,3,7,15]:
77         plt.figure(figsize=(8,4))
78         plt.plot(xnew, t_pred, label='Prediction')
79         plt.fill_between(xnew, mnew.ravel() + uncertainty, mnew.ravel() - uncertainty, alpha=0.1)
80         plt.plot(xnew, f(xnew), '--', color='black', label="Ground Truth", alpha=0.8)
81         plt.scatter(xnew, t_observed, color='orange', marker='.', linewidths=0.5, alpha=0.5, label="f(x) + noise")
82         plt.scatter(x, t, color='red', marker='x', linewidths=5, label='Samples')
83         plt.xlabel('x')
84         plt.title("Iteration #" + str(i+1))
85         plt.legend()
86         plt.show()

```

Listing 9: Gaussian Process Regression.